

Object detection and segmentation in deep learning


Colin Decourt – ANITI / NXP Semiconductors

08/02/2023

ANITI

Université
Fédérale
Toulouse
Midi-Pyrénées



From classification...



Class scores:
Cat: 0.85
Owl: 0.1
Dog: 0.03
...


Fully connected

... to object detection and segmentation




Semantic segmentation: what is it?

- Assign each pixel in the image to a category label (grass, cat, tree, sky...)
- Paired training data: for each training image, each pixel is labelled with a *semantic* category
- At test time, classify each pixel of a new image






Semantic segmentation idea: sliding window

- Classify many crops of the image using CNN (ResNet, VGG...)
- How to classify without context?
- How do we include context?



Extract patches

Classify centre pixel with CNN



Cat


Cat

Table

Semantic segmentation idea: sliding window




Very inefficient! For a 800x600 images ~ **58M boxes**
Not reusing shared features between overlapping patches

- Classify many crops of the image using CNN (ResNet, VGG...)
- How to classify without context?
- How do we include context?



Extract patches

Classify centre pixel with CNN




Cat

Cat

Table


Semantic segmentation idea: convolution

- Idea: encode image features with ConvNets, and perform semantic segmentation on top
- Classification architectures reduce spatial sizes to go deeper, but semantic segmentation requires the output size to be the same as the input size



Semantic segmentation idea: fully convolutional


Design a network with only convolutional layers without downsampling operators to make predictions for pixels all at ones



Very expensive as we don't decrease input size!


Semantic segmentation idea: fully convolutional

Design a network with only convolutional layers without downsampling operators to make predictions for pixels all at ones



Semantic segmentation idea: fully convolutional

- Design the network as a bunch of convolutional layers, with **downsampling** and **upsampling** inside the network!
- Downsampling can be pooling, strided convolution...
- Upsampling can be nearest neighbour interpolation, unpooling, transposed convolution




Semantic segmentation idea: fully convolutional

Idea: Learn to project low-resolution features onto pixel level (high-resolution)


- Design the network as a bunch of convolutional layers, with **downsampling** and **upsampling** inside the network!
- Downsampling can be pooling, strided convolution...
- Upsampling can be nearest neighbour interpolation, unpooling, transposed convolution



Input:
 $3 \times H \times W$




High-res:
 $D_1 \times H/2 \times W/2$




Med-res:
 $D_2 \times H/4 \times W/4$

Low-res:
 $D_3 \times H/4 \times W/4$



Med-res:
 $D_2 \times H/4 \times W/4$


High-res:
 $D_2 \times H/2 \times W/2$




Predictions:
 $H \times W$

Upsampling low-resolution features

Restore the condensed feature map to the original size of the input image by expanding the feature dimension



1. Unpooling




2. Transposed convolution


Transposed convolution

Learn to upsample the input feature map to the desired size using **learnable parameters**

Classic convolution





Transposed convolution




Transposed convolution

Learn to upsample the input feature map to the desired size using **learnable parameters**



Semantic segmentation: fully convolutional


- Design the network as a bunch of convolutional layers, with **downsampling** and **upsampling** inside the network!
- Downsampling can be pooling, strided convolution...
- Upsampling can be nearest neighbour interpolation, unpooling, transposed convolution





Don't differentiate instances (yet),
only care about pixels

Semantic segmentation: fully convolutional


- Design the network as a bunch of convolutional layers, with **downsampling** and **upsampling** inside the network!
- Downsampling can be pooling, strided convolution...
- Upsampling can be nearest neighbour interpolation, unpooling, transposed convolution




Semantic segmentation: some models...



From classification to object detection and segmentation



Object Detection



Instance Segmentation




Multiple Object

[This image](#) is [CC0 public domain](#)


Object detection: what is it?

- The task of assigning a label and a bounding box (or a mask) to all objects in the image
- Semantic segmentation:
 - No objects, just pixels
 - Output a segmentation mask with the same size as the input
- Object detection: for each object output **a label** and a box (or mask)




An easy case... Detecting single objects

Object detection is at the same time a **regression** problem and a **classification** task






Detecting single objects: multitask loss




Object detection: challenges

- Multiple outputs: **variable number of objects per image**
- Multiple types of output: category label, bounding boxes (orientation, velocity...)





Object detection: challenges

- Multiple outputs: **variable number of objects per image**
- Multiple types of output: category label, bounding boxes (orientation, velocity...)




Cat + (x, y, w, h)

4 numbers to predict




Object detection: challenges

- Multiple outputs: **variable number of objects per image**
- Multiple types of output: category label, bounding boxes (orientation, velocity...)



Cat + (x, y, w, h)


4 numbers to predict



Dog + (x, y, w, h)


12 numbers to predict

Dog + (x, y, w, h)
Cat + (x, y, w, h)




Object detection: challenges

- Multiple outputs: **variable number of objects per image**
- Multiple types of output: category label, bounding boxes (orientation, velocity...)



Cat + (x, y, w, h)

4 numbers to predict




Dog + (x, y, w, h)

Dog + (x, y, w, h)

Cat + (x, y, w, h)

12 numbers to predict



Duck + (x, y, w, h)

Duck + (x, y, w, h)

...

Many numbers to predict!!


Remember... The sliding window idea

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background

Very inefficient! For a 800x600 images ~ **58M boxes**
Need to apply CNN to a huge number of locations!



Extract patches



Cat? Yes
Table? No
Background? No


Remember... The sliding window idea

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background

Very inefficient! For a 800x600 images ~ **58M boxes**
Need to apply CNN to a huge number of locations!



Extract patches




Cat? Yes
Table? No
Background? No


Remember... The sliding window idea

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background

Very inefficient! For a 800x600 images ~ **58M boxes**
Need to apply CNN to a huge number of locations!



Extract patches




Cat? Yes
Table? No
Background? No

Remember... The sliding window idea

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background

Very inefficient! For a 800x600 images ~ **58M boxes**
Need to apply CNN to a huge number of locations!



Cat? No
Table? Yes
Background? No

Object detection: region proposals

Selective search algorithm :

1. Generate initial sub-segmentation, many candidate regions generation
2. Use greedy algorithm to recursively combine similar region into larger ones
 1. From set of regions, choose two that are most similar.
 2. Combine them into a single, larger region.
 3. Repeat the above steps for multiple iterations.
3. Use the generated regions to produce the final candidate region proposals




Object detection: R-CNN (Girshick et al., 2013)

- Classify EACH proposed region (SVM)
- Bounding box regression:
 - Predict “transform” to correct the proposed RoI
 - 4 numbers: (t_x, t_y, t_h, t_w)
- Final output:
 - Proposal: (p_x, p_y, p_h, p_w)
 - Transform: (t_x, t_y, t_h, t_w) (a “*correction*”)
 - Output box: (b_x, b_y, b_h, b_w)
 - $b_x = p_x + p_w t_x$ and $b_y = p_y + p_h t_y$
 - $b_w = p_w e^{t_w}$ and $b_h = p_h e^{t_h}$




Object detection: R-CNN (Girshick et al., 2013)

- Classify EACH proposed region (SVM)
- Bounding box regression:
 - Predict “transform” to correct the proposed RoI
 - 4 numbers: (t_x, t_y, t_h, t_w)
- Final output:
 - Proposal: (p_x, p_y, p_h, p_w)
 - Transform: (t_x, t_y, t_h, t_w) (a “*correction*”)
 - Output box: (b_x, b_y, b_h, b_w)
 - $b_x = p_x + p_w t_x$ and $b_y = p_y + p_h t_y$
 - $b_w = p_w e^{t_w}$ and $b_h = p_h e^{t_h}$




Object detection: R-CNN (Girshick et al., 2013)

- Classify EACH proposed region (SVM)
- Bounding box regression:
 - Predict “transform” to correct the proposed RoI
 - 4 numbers: (t_x, t_y, t_h, t_w)
- Final output:
 - Proposal: (p_x, p_y, p_h, p_w)
 - Transform: (t_x, t_y, t_h, t_w) (a “*correction*”)
 - Output box: (b_x, b_y, b_h, b_w)
 - $b_x = p_x + p_w t_x$ and $b_y = p_y + p_h t_y$
 - $b_w = p_w e^{t_w}$ and $b_h = p_h e^{t_h}$




Object detection: R-CNN (Girshick et al., 2013)

- Classify EACH proposed region (SVM)
- Bounding box regression:
 - Predict “transform” to correct the proposed RoI
 - 4 numbers: (t_x, t_y, t_h, t_w)
- Final output:
 - Proposal: (p_x, p_y, p_h, p_w)
 - Transform: (t_x, t_y, t_h, t_w) (a “*correction*”)
 - Output box: (b_x, b_y, b_h, b_w)
 - $b_x = p_x + p_w t_x$ and $b_y = p_y + p_h t_y$
 - $b_w = p_w e^{t_w}$ and $b_h = p_h e^{t_h}$




Object detection: R-CNN (Girshick et al., 2013)

- Classify EACH proposed region (SVM)
- Bounding box regression:
 - Predict “transform” to correct the proposed RoI
 - 4 numbers: (t_x, t_y, t_h, t_w)
- Final output:
 - Proposal: (p_x, p_y, p_h, p_w)
 - Transform: (t_x, t_y, t_h, t_w) (a “*correction*”)
 - Output box: (b_x, b_y, b_h, b_w)
 - $b_x = p_x + p_w t_x$ and $b_y = p_y + p_h t_y$
 - $b_w = p_w e^{t_w}$ and $b_h = p_h e^{t_h}$



Object detection: R-CNN (Girshick et al., 2013)


- Classify EACH proposed region (SVM)
- Bounding box regression:
 - Predict “transform” to correct the proposed RoI
 - 4 numbers: (t_x, t_y, t_h, t_w)
- Final output:
 - Proposal: (p_x, p_y, p_h, p_w)
 - Transform: (t_x, t_y, t_h, t_w) (a “*correction*”)
 - Output box: (b_x, b_y, b_h, b_w)
 - $b_x = p_x + p_w t_x$ and $b_y = p_y + p_h t_y$
 - $b_w = p_w e^{t_w}$ and $b_h = p_h e^{t_h}$



Object detection: R-CNN (Girshick et al., 2013)

- Classify EACH proposed region (SVM)
- Bounding box regression:
 - Predict “transform” to correct the proposed RoI
 - 4 numbers: (t_x, t_y, t_h, t_w)
- Final output:
 - Proposal: (p_x, p_y, p_h, p_w)
 - Transform: (t_x, t_y, t_h, t_w) (a “*correction*”)
 - Output box: (b_x, b_y, b_h, b_w)
 - $b_x = p_x + p_w t_x$ and $b_y = p_y + p_h t_y$
 - $b_w = p_w e^{t_w}$ and $b_h = p_h e^{t_h}$


Problem: very slow! Need 2000 independent forward pass for each image



Object detection: “Slow” R-CNN (Girshick et al., 2013)

Problem: very slow! Need 2000 independent forward pass for each image

- Classify EACH proposed region (SVM)
- Bounding box regression:
 - Predict “transform” to correct the proposed RoI
 - 4 numbers: (t_x, t_y, t_h, t_w)
- Final output:
 - Proposal: (p_x, p_y, p_h, p_w)
 - Transform: (t_x, t_y, t_h, t_w) (a “*correction*”)
 - Output box: (b_x, b_y, b_h, b_w)
 - $b_x = p_x + p_w t_x$ and $b_y = p_y + p_h t_y$
 - $b_w = p_w e^{t_w}$ and $b_h = p_h e^{t_h}$



Object detection: Fast-RCNN (Girshick et al., 2015)




Object detection: Fast-RCNN (Girshick et al., 2015)




Most of computations
happens here (ResNet,
VGG, AlexNet)

Pre-trained
ConvNet


Object detection: Fast-RCNN (Girshick et al., 2015)




Object detection: Fast-RCNN (Girshick et al., 2015)




Object detection: Fast-RCNN (Girshick et al., 2015)




Object detection: Fast-RCNN (Girshick et al., 2015)




Object detection: Fast-RCNN (Girshick et al., 2015)



Object detection: Fast-RCNN (Girshick et al., 2015)



Object detection: Fast-RCNN (Girshick et al., 2015)




RoI pooling




Input image (e.g.
 $3 \times 640 \times 480$)



Feature map
(e.g. $512 \times 20 \times 15$)




RoI pooling




Input image (e.g.
 $3 \times 640 \times 480$)




Feature map
(e.g. $512 \times 20 \times 15$)




RoI pooling




Input image (e.g.
 $3 \times 640 \times 480$)




Roi pooling




Input image (e.g.
 $3 \times 640 \times 480$)




ROI pooling




Input image (e.g.
 $3 \times 640 \times 480$)




Roi pooling




Input image (e.g.
 $3 \times 640 \times 480$)



RoI pooling




Input image (e.g.
 $3 \times 640 \times 480$)




Problem: region features might be slightly misaligned → RoI align

RoI pooling




Input image (e.g.
 $3 \times 640 \times 480$)




Object detection: Fast R-CNN vs. “Slow” R-CNN

Runtime dominated by region proposals



Girshick et al, “Rich feature hierarchies for accurate object detection and semantic segmentation”, CVPR 2014.
He et al, “Spatial pyramid pooling in deep convolutional networks for visual recognition”, ECCV 2014
Girshick, “Fast R-CNN”, ICCV 2015

Object detection: learn to *propose* regions, Faster R-CNN (Ren et al., 2015)




- Idea: insert a **Region Proposal Network (RPN)** to predict proposals from features
- Otherwise same as Fast R-CNN: crop features for each proposal, classify each one

Object detection: Region Proposal Network (RPN)

Imagine an **anchor box** of fixed size at each point in the feature map



Input image (e.g.
 $3 \times 640 \times 480$)




Object detection: Region Proposal Network (RPN)

Imagine an anchor box of fixed size at each point in the feature map



Input image (e.g.
3x640x480)




Project proposal onto
features

Object detection: Region Proposal Network (RPN)

Imagine an **anchor box** of fixed size at each point in the feature map



Input image (e.g.
3x640x480)




Object detection: Region Proposal Network (RPN)

Imagine an **anchor box** of fixed size at each point in the feature map




Input image (e.g.
 $3 \times 640 \times 480$)




Object detection: Region Proposal Network (RPN)



Input image (e.g.
 $3 \times 640 \times 480$)




Object detection: learn to *propose* regions, Faster R-CNN (Ren et al., 2015)



- Jointly train with 4 losses:
 1. **RPN classification:** anchor box is object / not an object
 2. **RPN regression:** predict transform from anchor box to proposal box
 3. **Object classification:** classify proposals as background / object class
 4. **Object regression:** predict transform from proposal to object box


Object detection: learn to *propose* regions, Faster R-CNN (Ren et al., 2015)



- Jointly train with 4 losses:
 1. **RPN classification:** anchor box is object / not an object
 2. **RPN regression:** predict transform from anchor box to proposal box
 3. **Object classification:** classify proposals as background / object class
 4. **Object regression:** predict transform from proposal to object box

**Test time speed: 0.2 seconds vs 2.3 seconds
(Fast R-CNN) vs 49 seconds (R-CNN)!**

Object detection: learn to *propose* regions, Faster R-CNN (Ren et al., 2015)



- Faster R-CNN is a two-stage object detector:
 1. Extract features using a backbone network and propose regions (mostly background)
 2. For each region: crop features, predict object class and bounding box offset
- Single-stage object detectors: YOLO, SSD, RetinaNet


Object detection: single-stage object detectors (YOLO, SSD, RetinaNet)

- Predict object class and location in ONE single step
- Similar to RPN of Faster R-CNN
- Predict the position of the box AND the class of the object in a box

Object detection: YOLO (Redmond et al., 2015)

Object detection: YOLO (Redmond et al., 2015)


1. Divide the image into cells with an $S \times S$ grid



$S=7$

Object detection: YOLO (Redmond et al., 2015)


2. Each cell predicts B bounding boxes



Same idea for YOLO v2, v3, v4, ..., v8!


Object detection: YOLO (Redmond et al., 2015)

3. Return bounding boxes above confidence threshold





All other bounding boxes have a confidence probability less than the threshold (e.g 0.9) so they are suppressed.

- Each cell predicts:
 - For each anchor box:
 - 4 (box offset) coordinates (dx, dy, dh, dw)
 - 1 confidence value
 - Class probabilities (80 for COCO dataset, 20 for PASCAL-VOC)
- Output: $7 \times 7 \times (5 * B + C)$
- Similar to RPN!



Object detection: single stage object detectors



Object detection (and segmentation): but...

Object detection (and segmentation): but...

Object detection/segmentation is a wide field of research, and many architectures exist.

Object detection (and segmentation): but...

Object detection/segmentation is a wide field of research, and many architectures exist.

- Alternatives to anchor-based methods exist:
 - CenterNet (Duan et al., 2019)
 - FCOS (Tian et al., 2019)
 - R-FCN (Dai et al., 2016)
 - DETR (object detection with transformers) (Carion et al., 2020)

Object detection (and segmentation): but...

Object detection/segmentation is a wide field of research, and many architectures exist.


- Alternatives to anchor-based methods exist:
 - CenterNet (Duan et al., 2019)
 - FCOS (Tian et al., 2019)
 - R-FCN (Dai et al., 2016)
 - DETR (object detection with transformers) (Carion et al., 2020)
- Lot of variables:
 - Backbone network: VGG, ResNet, InceptionV2/V3, MobileNet, EfficientNet
 - Architecture style: two-stage, single-stage, hybrid...

Object detection (and segmentation): but...

Object detection/segmentation is a wide field of research, and many architectures exist.

- Alternatives to anchor-based methods exist:
 - CenterNet (Duan et al., 2019)
 - FCOS (Tian et al., 2019)
 - R-FCN (Dai et al., 2016)
 - DETR (object detection with transformers) (Carion et al., 2020)
- Lot of variables:
 - Backbone network: VGG, ResNet, InceptionV2/V3, MobileNet, EfficientNet
 - Architecture style: two-stage, single-stage, hybrid...
- **Takeways:**
 - Two-stage detectors are slower but more accurate
 - Single-stage detectors are faster but not as accurate
 - Bigger / Deeper backbones work better


Instance segmentation: Mask R-CNN



Add a small mask network that operates on each RoI and predicts 28x28 binary mask


Instance segmentation: Mask R-CNN

Add a small mask network that operates on each RoI and predicts 28×28 binary mask



Object detection and segmentation datasets

- Pascal VOC dataset:
 - Detection, classification, segmentation
 - 10000 images with 20 categories
- COCO dataset:
 - Caption generation, object detection, key point detection and object segmentation
 - 120000 images for training / 40000 for validation with 80 categories
- KITTI autonomous driving dataset:
 - Detection, classification, semantic and instance segmentation, tracking...



Object detection and segmentation: frameworks

Many implementations of the aforementioned model are available on GitHub.

TensorFlow Object Detection API:

https://github.com/tensorflow/models/tree/master/research/object_detection

Detectron2 (PyTorch): <https://github.com/facebookresearch/detectron2>

Torchvision (PyTorch): <https://pytorch.org/vision/stable/index.html>

Use pre-trained model to finetune on your own dataset!

Other object detection/segmentation tasks


3D object detection



Key point object detection



3D semantic segmentation






Appendix

ANITI

Université
Fédérale
Toulouse
Midi-Pyrénées

Object detection (or segmentation): comparing boxes (or masks)


- How can we compare the prediction and the bounding boxes?



Object detection (or segmentation): comparing boxes (or masks)


- How can we compare the prediction and the bounding boxes?
- Use the **Intersection over Union (IoU)**:

$$IoU = \frac{\text{Area of Intersection}}{\text{Area of Union}}$$



Object detection (or segmentation): comparing boxes (or masks)

- How can we compare the prediction and the bounding boxes?
 - Use the **Intersection over Union (IoU)**:
- $$IoU = \frac{\text{Area of Intersection}}{\text{Area of Union}}$$
- IoU > 0.5 is “decent”




Object detection (or segmentation): comparing boxes (or masks)

- How can we compare the prediction and the bounding boxes?

- Use the **Intersection over Union (IoU)**:

$$IoU = \frac{\text{Area of Insersection}}{\text{Area of Union}}$$

- IoU > 0.5 is “decent”
- IoU > 0.7 is “pretty good”




Object detection (or segmentation): comparing boxes (or masks)

- How can we compare the prediction and the bounding boxes?
- Use the **Intersection over Union (IoU)**:


$$IoU = \frac{\text{Area of Insersection}}{\text{Area of Union}}$$

- IoU > 0.5 is “decent”
- IoU > 0.7 is “pretty good”
- IoU > 0.9 is “almost perfect”
- For segmentation masks this is done pixel-wise



Object detection: deal with overlapping boxes

- **Problem:** object detectors often output many overlapping detection (due to multiple anchors per pixel)
- **Solution:** post-process raw detections using **Non-Max Suppression (NMS)**
- Algorithm:
 1. Select highest-scoring box
 2. Eliminate lower-scoring boxes with $\text{IoU} > \text{threshold}$ (e.g. 0.7)
 3. If any boxes remain, GOTO 1



Puppy image is CC0 Public Domain


Object detection: deal with overlapping boxes

- **Problem:** object detectors often output many overlapping detection (due to multiple anchors per pixel)
- **Solution:** post-process raw detections using **Non-Max Suppression (NMS)**
- Algorithm:
 1. Select highest-scoring box
 2. Eliminate lower-scoring boxes with $\text{IoU} > \text{threshold}$ (e.g. 0.7)
 3. If any boxes remain, GOTO 1

$$\text{IoU}(\text{blue}, \text{orange}) = 0.78$$

$$\text{IoU}(\text{blue}, \text{purple}) = 0.05$$

$$\text{IoU}(\text{blue}, \text{yellow}) = 0.07$$




Puppy image is CC0 Public Domain

Object detection: deal with overlapping boxes

- **Problem:** object detectors often output many overlapping detection (due to multiple anchors per pixel)
- **Solution:** post-process raw detections using **Non-Max Suppression (NMS)**
- Algorithm:
 1. Select highest-scoring box
 2. Eliminate lower-scoring boxes with $\text{IoU} > \text{threshold}$ (e.g. 0.7)
 3. If any boxes remain, GOTO 1


$$\text{IoU}(\text{purple box}, \text{yellow box}) = 0.74$$



Puppy image is CC0 Public Domain

Object detection: deal with overlapping boxes


- **Problem:** object detectors often output many overlapping detection (due to multiple anchors per pixel)
- **Solution:** post-process raw detections using **Non-Max Suppression (NMS)**
- Algorithm:
 1. Select highest-scoring box
 2. Eliminate lower-scoring boxes with $\text{IoU} > \text{threshold}$ (e.g. 0.7)
 3. If any boxes remain, GOTO 1
- **Problem:** NMS may eliminate “good” boxes when objects are highly overlapping...



Puppy image is CC0 Public Domain


Evaluating object detectors – Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = Area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)




Evaluating object detectors – Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = Area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with $\text{IoU} > 0.5$, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative




Evaluating object detectors – Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = Area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with $\text{IoU} > 0.5$, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve




Evaluating object detectors – Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = Area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with $\text{IoU} > 0.5$, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve




Evaluating object detectors – Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = Area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with IoU > 0.5, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve




Evaluating object detectors – Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = Area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with IoU > 0.5, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve



Evaluating object detectors – Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = Area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with IoU > 0.5, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve



Evaluating object detectors – Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
 2. For each category, compute Average Precision (AP) = Area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with IoU > 0.5, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve
 2. Average Precision (AP) = Area under PR curve
 3. Mean Average Precision (mAP) = average of AP for each category
- CarAP = 0.65
Cat AP = 0.80
Dog AP = 0.86
mAP@0.5 = 0.77

Evaluating object detectors – Mean Average Precision (mAP)

1. Run object detector on all test images (with NMS)
 2. For each category, compute Average Precision (AP) = Area under Precision vs Recall Curve
 1. For each detection (highest score to lowest score)
 1. If it matches some GT box with $\text{IoU} > 0.5$, mark it as positive and eliminate the GT
 2. Otherwise mark it as negative
 3. Plot a point on PR Curve
 2. Average Precision (AP) = Area under PR curve
 3. Mean Average Precision (mAP) = average of AP for each category
 4. For “COCO mAP”: Compute mAP@thresh for each IoU threshold (0.5, 0.55, 0.6, ..., 0.95) and take average
- mAP@0.50 = 0.77
mAP@0.55 = 0.71
mAP@0.60 = 0.65
...
mAP@0.95 = 0.2
- COCO mAP = 0.4